

fina_project

May 14, 2020

Created on Tue Apr 21 11:10:11 2020

@author: Deeps

```
[1]: import warnings
warnings.filterwarnings("ignore")
#!pip install torch
#!pip install torchsummary
#!pip install torchvision
```

```
[2]: import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
from tqdm import tqdm # Displays a progress bar

import torch
from torch.utils import data
from torchsummary import summary
from torch import nn
from torch import optim
from PIL import Image, ImageFilter
from sklearn.model_selection import train_test_split as tts
from sklearn.metrics import roc_auc_score, accuracy_score, confusion_matrix
import torch.nn.functional as F
from torchvision import datasets, models, transforms
from torch.utils.data import Dataset, Subset, DataLoader, random_split
from data_gen import Dataset
from augment import Augmentations as aug
from visualization import visualizer
```

0.1 Load the dataset and train, val, test splits

```
[3]: num_epochs = 10
num_classes = 2
batch_size = 5
num_rotations = 2 #180
rotations = [0,180]
```

```

learning_rate = 0.001
image_size = (748,500)
resnet_resize = (374,250)
datacsv = pd.read_csv("data.csv")
cv_results = []

```

0.1.1 —————Data Augmentation—————

Since we only have 119 images, and I want to keep the test set fixed at 20, with exactly 3 examples of class 0, I first split the data into 99 and 20. I then only apply augmentation to the 99 for training. I do not apply this augmentation function to the test data. only the transformations are applied.

```

[4]: all_ = len(datacsv['image'])
      train_size = int(all_ *(5/6))
      test_size = all_-train_size

      traincsv = datacsv[:train_size]
      testcsv = datacsv[train_size:]

[5]: data_base_aug = dict(traincsv)
      data_base_aug["image"] = list(data_base_aug["image"])
      data_base_aug["label"] = list(data_base_aug["label"])

      if num_rotations>0:
          rot = aug(rotations)
          for k in range(len(data_base_aug["image"])):
              result = rot.
      ↪rotate_append(data_base_aug["image"][k],data_base_aug["label"][k])
          traincsv = traincsv.append(result,ignore_index = True)

```

0.2 —————inspect—————

```

[14]: x_train = traincsv['image']
      x_test = testcsv['image']
      y_train = traincsv['label']
      y_test = testcsv['label']

      #x_train,x_test,y_train,y_test = tts(datacsv["image"],↵
      ↪datacsv["label"],test_size=1/6, shuffle=False, random_state=42)
      partition = {'train':list(x_train),'validation':list(x_test)}
      labels = {}

      combinedforlabels = traincsv.append(testcsv,ignore_index = True)

      for k in range(len(combinedforlabels['image'])):
          labels['%s'%combinedforlabels['image'][k]] = combinedforlabels['label'][k]

```

```

print("Total original Images: %s"%all_)
print('Train images( after augmentation): %s'%(len(y_train)), ', Test images:␣
→%s'%(len(y_test)))
print('Train Positives(with augmentation): %s'%(np.sum(y_train)/len(y_train)),␣
→', Test Positives(No augmentation): %s'%(np.sum(y_test)/len(y_test)))
print('Total Real dataset positives: %s'%(np.sum(datacsv["label"])/
→len(datacsv["label"])))

view_image = Image.open('Dataset/' + str(int(combinedforlabels["image"][88])) +␣
→'.jpg').convert('LA')
plt.imshow(view_image)

```


Total original Images: 119

Train images(after augmentation): 183 , Test images: 20

Train Positives(with augmentation): 0.819672131147541 , Test Positives(No
augmentation): 0.85

Total Real dataset positives: 0.8823529411764706

[14]: <matplotlib.image.AxesImage at 0x2ef2a968f08>



1 ———-Create dataset generator—————

```
[16]: for cv in range(1):
    #x_train,x_test = tts(datacsv["image"],test_size=1/6, shuffle=False)

    #####

    training_set = Dataset(partition["train"],labels, image_size)
    training_generator = data.DataLoader(training_set, batch_size=batch_size,
    ↪shuffle=True)

    validation_set = Dataset(partition['validation'], labels, image_size)
    validation_generator = data.DataLoader(validation_set,
    ↪batch_size=batch_size, shuffle=True)

    #####
    class Network(nn.Module):
        def __init__(self):
            super().__init__()
            self.resnet = torch.hub.load('pytorch/vision', 'resnet34',
    ↪pretrained=True)

            self.num_fters = self.resnet.fc.in_features
            self.resnet.fc = nn.Linear(self.num_fters,8)
            self.fc1 = nn.Linear(8, num_classes)

        def forward(self,x):
            # TODO: Design your own network, implement forward pass here
            x = F.dropout(F.relu(self.resnet(x)),0.2)#3*748*512 -> 6*744*508 ->
    ↪6*372*254
            out = F.softmax(self.fc1(x))
            return out

    device = "cuda" if torch.cuda.is_available() else "cpu" # Configure device
    model = Network().to(device)

    ct = 0
    model_size = len([1 for k in model.children()])
    print("model components: ", model_size)
    for child in model.children():
        ct+=1
        if ct==1:
            for param in child.parameters():
                param.requires_grad=False

    model.resnet.fc.weight.requires_grad = True
    model.resnet.fc.bias.requires_grad = True
```

```

#model.resnet.layer4[1].conv2.weight.requires_grad = True

model.fc1.weight.requires_grad = True
model.fc1.bias.requires_grad = True

weights = torch.tensor([0.6,0.4])
criterion = nn.CrossEntropyLoss(weight = weights) # Specify the loss layer
optimizer = optim.Adam(model.parameters()) # Specify optimizer and assign
↳ trainable parameters to it, weight_decay is L2 regularization strength

def train(model, training_generator, num_epoch = num_epochs): # Train the
↳ model
    print("Start training...")
    model.train() # Set the model to training mode
    train_val_loss = {'train': [], 'validation': []}
    for i in range(num_epoch):
        running_loss = []
        accuracy = []
        for batch, label in tqdm(training_generator):
            batch = batch.to(device)
            label = label.to(device)
            optimizer.zero_grad() # Clear gradients from the previous
↳ iteration

            pred = model(batch) # This will call Network.forward() that you
↳ implement

            loss = criterion(pred, label) # Calculate the loss
            running_loss.append(loss.item())
            correct = (torch.argmax(pred,dim=1)==label).sum().item()
            accuracy.append(correct/batch_size)

            loss.backward() # Backprop gradients to all tensors in the
↳ network

            optimizer.step() # Update trainable weights

        val_loss, val_accuracy = evaluate(model, validation_generator)
        print("Epoch {}  loss:{}  eval_loss:{}  accuracy:{}  "
↳ eval_accuracy: {}".format(i+1,np.mean(running_loss), val_loss, np.
↳ mean(accuracy),val_accuracy)) # Print the average loss for this epoch

        train_val_loss['train'].append(np.mean(running_loss))
        train_val_loss['validation'].append(np.mean(val_loss))

    if (i+1)%10==0 or i==(num_epoch-1) or val_accuracy>0.85:

```

```

        try:
            torch.save(model, "Trained_model/teeth_model_%s.pth"%(i+1))
            print("Model saved at Trained_model")
        except:
            print("Could not save model")
    if val_accuracy>0.85:
        print("early stop")
        break
print("Done!")
return train_val_loss

def evaluate(model, validation_generator): # Evaluate accuracy on
    validation / test set
    model.eval() # Set the model to evaluation mode
    correct = 0
    total = 0
    val_run_loss = []
    with torch.set_grad_enabled(True): # Do not calculate gradient to speed
    up computation
        for batch, label in tqdm(validation_generator):
            batch = batch.to(device)
            label = label.to(device)
            pred = model(batch)
            loss = criterion(pred, label).item()
            val_run_loss.append(loss)
            correct += (torch.argmax(pred,dim=1)==label).sum().item()
            total+=batch_size
    acc = correct/total
    #print("Evaluation accuracy: {}".format(acc))
    return (np.mean(loss),acc)

def predict(model, validation_generator): # Evaluate accuracy on validation
    / test set
    model.eval() # Set the model to evaluation mode
    results = {'pred':[], 'real':[]}
    with torch.set_grad_enabled(True): # Do not calculate gradient to speed
    up computation
        for batch, label in tqdm(validation_generator):
            batch = batch.to(device)
            label = label.to(device)
            pred = model(batch)
            pred = torch.argmax(pred,dim=1)
            for k in range(len(pred)):
                results['pred'].append(pred[k].item())
                results['real'].append(label[k].item())
    return results

```

Using cache found in C:\Users\Deeps\.cache\torch\hub\pytorch_vision_master
model components: 2

1.1 train and eval

```
[17]: train_history = train(model, training_generator, num_epochs)
```

Start training...

```
100%|  
  | 37/37 [01:42<00:00, 2.76s/it]
```

```
100%|  
  | 4/4 [00:10<00:00, 2.59s/it]
```

```
Epoch 1  loss:0.6036993192659842  eval_loss:0.3954412639141083  
accuracy:0.7837837837837839  eval_accuracy: 0.85
```

```
100%|  
  | 37/37 [01:38<00:00, 2.65s/it]
```

```
100%|  
  | 4/4 [00:11<00:00, 2.80s/it]
```

```
Epoch 2  loss:0.5651485364179354  eval_loss:0.5955150127410889  
accuracy:0.8108108108108109  eval_accuracy: 0.85
```

```
100%|  
  | 37/37 [01:58<00:00, 3.21s/it]
```

```
100%|  
  | 4/4 [00:13<00:00, 3.35s/it]
```

```
Epoch 3  loss:0.5517468508836385  eval_loss:0.3136637508869171  
accuracy:0.8108108108108109  eval_accuracy: 0.85
```

```
100%|  
  | 37/37 [02:07<00:00, 3.45s/it]
```

```
100%|  
  | 4/4 [00:16<00:00, 4.22s/it]
```

```
Epoch 4  loss:0.5464425755513681  eval_loss:0.587364673614502  
accuracy:0.810810810810811  eval_accuracy: 0.85
```

```
100%|  
  | 37/37 [02:19<00:00, 3.76s/it]
```

```
100%|  
  | 4/4 [00:13<00:00, 3.49s/it]
```

```
Epoch 5  loss:0.5428367458485268  eval_loss:0.37417370080947876  
accuracy:0.810810810810811  eval_accuracy: 0.85
```

```
100%|  
  | 37/37 [02:19<00:00, 3.78s/it]
```

```
100%|  
  | 4/4 [00:15<00:00, 3.89s/it]
```

```

Epoch 6   loss:0.5416768104643435   eval_loss:0.45081761479377747
accuracy:0.8108108108108109   eval_accuracy: 0.85
100%|
  | 37/37 [02:18<00:00,  3.76s/it]
100%|
  | 4/4 [00:15<00:00,  3.92s/it]

Epoch 7   loss:0.5464279063650079   eval_loss:0.5864800214767456
accuracy:0.8108108108108109   eval_accuracy: 0.85
100%|
  | 37/37 [02:16<00:00,  3.69s/it]
100%|
  | 4/4 [00:15<00:00,  3.94s/it]

Epoch 8   loss:0.5505075970211545   eval_loss:0.5872970819473267
accuracy:0.8108108108108109   eval_accuracy: 0.85
100%|
  | 37/37 [02:25<00:00,  3.92s/it]
100%|
  | 4/4 [00:13<00:00,  3.46s/it]


Epoch 9   loss:0.5554544079948116   eval_loss:0.3721107244491577
accuracy:0.8108108108108109   eval_accuracy: 0.85
100%|
  | 37/37 [02:22<00:00,  3.86s/it]
100%|
  | 4/4 [00:13<00:00,  3.39s/it]

Epoch 10   loss:0.5490918199758272   eval_loss:0.5857167840003967
accuracy:0.8108108108108109   eval_accuracy: 0.85
Model saved at Trained_model
Done!

```

1.2 Evaluation

```
[18]: kk = pd.DataFrame(train_history).plot(kind='line',grid=True)
```

1.3 See Confusion matrix

```
[19]: see_pred = pd.DataFrame(predict(model, validation_generator))
      print("Confusion matrix shows the TP, FP,TN, FN rates of the model")
      print(confusion_matrix(see_pred['real'],see_pred['pred']))
```


100%|

| 4/4 [00:09<00:00, 2.33s/it]

Confusion matrix shows the TP, FP,TN, FN rates of the model

```
[[ 0  3]
 [ 0 17]]
```

```
[20]: vis=visualizer()
      vis.plot_filter(layer=model.resnet.conv1,single_channel = True)
```



2 ———The END—————

```
[21]: summary(model,(3,resnet_resize[0],resnet_resize[1]))
```

Layer (type)	Output Shape	Param #

Conv2d-1	[-1, 64, 187, 125]	9,408
BatchNorm2d-2	[-1, 64, 187, 125]	128
ReLU-3	[-1, 64, 187, 125]	0
MaxPool2d-4	[-1, 64, 94, 63]	0
Conv2d-5	[-1, 64, 94, 63]	36,864
BatchNorm2d-6	[-1, 64, 94, 63]	128
ReLU-7	[-1, 64, 94, 63]	0
Conv2d-8	[-1, 64, 94, 63]	36,864
BatchNorm2d-9	[-1, 64, 94, 63]	128
ReLU-10	[-1, 64, 94, 63]	0
BasicBlock-11	[-1, 64, 94, 63]	0
Conv2d-12	[-1, 64, 94, 63]	36,864
BatchNorm2d-13	[-1, 64, 94, 63]	128
ReLU-14	[-1, 64, 94, 63]	0
Conv2d-15	[-1, 64, 94, 63]	36,864
BatchNorm2d-16	[-1, 64, 94, 63]	128
ReLU-17	[-1, 64, 94, 63]	0
BasicBlock-18	[-1, 64, 94, 63]	0
Conv2d-19	[-1, 64, 94, 63]	36,864
BatchNorm2d-20	[-1, 64, 94, 63]	128
ReLU-21	[-1, 64, 94, 63]	0
Conv2d-22	[-1, 64, 94, 63]	36,864
BatchNorm2d-23	[-1, 64, 94, 63]	128
ReLU-24	[-1, 64, 94, 63]	0
BasicBlock-25	[-1, 64, 94, 63]	0
Conv2d-26	[-1, 128, 47, 32]	73,728
BatchNorm2d-27	[-1, 128, 47, 32]	256
ReLU-28	[-1, 128, 47, 32]	0
Conv2d-29	[-1, 128, 47, 32]	147,456
BatchNorm2d-30	[-1, 128, 47, 32]	256
Conv2d-31	[-1, 128, 47, 32]	8,192
BatchNorm2d-32	[-1, 128, 47, 32]	256
ReLU-33	[-1, 128, 47, 32]	0
BasicBlock-34	[-1, 128, 47, 32]	0
Conv2d-35	[-1, 128, 47, 32]	147,456
BatchNorm2d-36	[-1, 128, 47, 32]	256
ReLU-37	[-1, 128, 47, 32]	0
Conv2d-38	[-1, 128, 47, 32]	147,456
BatchNorm2d-39	[-1, 128, 47, 32]	256
ReLU-40	[-1, 128, 47, 32]	0
BasicBlock-41	[-1, 128, 47, 32]	0

Conv2d-42	[-1, 128, 47, 32]	147,456
BatchNorm2d-43	[-1, 128, 47, 32]	256
ReLU-44	[-1, 128, 47, 32]	0
Conv2d-45	[-1, 128, 47, 32]	147,456
BatchNorm2d-46	[-1, 128, 47, 32]	256
ReLU-47	[-1, 128, 47, 32]	0
BasicBlock-48	[-1, 128, 47, 32]	0
Conv2d-49	[-1, 128, 47, 32]	147,456
BatchNorm2d-50	[-1, 128, 47, 32]	256
ReLU-51	[-1, 128, 47, 32]	0
Conv2d-52	[-1, 128, 47, 32]	147,456
BatchNorm2d-53	[-1, 128, 47, 32]	256
ReLU-54	[-1, 128, 47, 32]	0
BasicBlock-55	[-1, 128, 47, 32]	0
Conv2d-56	[-1, 256, 24, 16]	294,912
BatchNorm2d-57	[-1, 256, 24, 16]	512
ReLU-58	[-1, 256, 24, 16]	0
Conv2d-59	[-1, 256, 24, 16]	589,824
BatchNorm2d-60	[-1, 256, 24, 16]	512
Conv2d-61	[-1, 256, 24, 16]	32,768
BatchNorm2d-62	[-1, 256, 24, 16]	512
ReLU-63	[-1, 256, 24, 16]	0
BasicBlock-64	[-1, 256, 24, 16]	0
Conv2d-65	[-1, 256, 24, 16]	589,824
BatchNorm2d-66	[-1, 256, 24, 16]	512
ReLU-67	[-1, 256, 24, 16]	0
Conv2d-68	[-1, 256, 24, 16]	589,824
BatchNorm2d-69	[-1, 256, 24, 16]	512
ReLU-70	[-1, 256, 24, 16]	0
BasicBlock-71	[-1, 256, 24, 16]	0
Conv2d-72	[-1, 256, 24, 16]	589,824
BatchNorm2d-73	[-1, 256, 24, 16]	512
ReLU-74	[-1, 256, 24, 16]	0
Conv2d-75	[-1, 256, 24, 16]	589,824
BatchNorm2d-76	[-1, 256, 24, 16]	512
ReLU-77	[-1, 256, 24, 16]	0
BasicBlock-78	[-1, 256, 24, 16]	0
Conv2d-79	[-1, 256, 24, 16]	589,824
BatchNorm2d-80	[-1, 256, 24, 16]	512
ReLU-81	[-1, 256, 24, 16]	0
Conv2d-82	[-1, 256, 24, 16]	589,824
BatchNorm2d-83	[-1, 256, 24, 16]	512
ReLU-84	[-1, 256, 24, 16]	0
BasicBlock-85	[-1, 256, 24, 16]	0
Conv2d-86	[-1, 256, 24, 16]	589,824
BatchNorm2d-87	[-1, 256, 24, 16]	512
ReLU-88	[-1, 256, 24, 16]	0
Conv2d-89	[-1, 256, 24, 16]	589,824

BatchNorm2d-90	[-1, 256, 24, 16]	512
ReLU-91	[-1, 256, 24, 16]	0
BasicBlock-92	[-1, 256, 24, 16]	0
Conv2d-93	[-1, 256, 24, 16]	589,824
BatchNorm2d-94	[-1, 256, 24, 16]	512
ReLU-95	[-1, 256, 24, 16]	0
Conv2d-96	[-1, 256, 24, 16]	589,824
BatchNorm2d-97	[-1, 256, 24, 16]	512
ReLU-98	[-1, 256, 24, 16]	0
BasicBlock-99	[-1, 256, 24, 16]	0
Conv2d-100	[-1, 512, 12, 8]	1,179,648
BatchNorm2d-101	[-1, 512, 12, 8]	1,024
ReLU-102	[-1, 512, 12, 8]	0
Conv2d-103	[-1, 512, 12, 8]	2,359,296
BatchNorm2d-104	[-1, 512, 12, 8]	1,024
Conv2d-105	[-1, 512, 12, 8]	131,072
BatchNorm2d-106	[-1, 512, 12, 8]	1,024
ReLU-107	[-1, 512, 12, 8]	0
BasicBlock-108	[-1, 512, 12, 8]	0
Conv2d-109	[-1, 512, 12, 8]	2,359,296
BatchNorm2d-110	[-1, 512, 12, 8]	1,024
ReLU-111	[-1, 512, 12, 8]	0
Conv2d-112	[-1, 512, 12, 8]	2,359,296
BatchNorm2d-113	[-1, 512, 12, 8]	1,024
ReLU-114	[-1, 512, 12, 8]	0
BasicBlock-115	[-1, 512, 12, 8]	0
Conv2d-116	[-1, 512, 12, 8]	2,359,296
BatchNorm2d-117	[-1, 512, 12, 8]	1,024
ReLU-118	[-1, 512, 12, 8]	0
Conv2d-119	[-1, 512, 12, 8]	2,359,296
BatchNorm2d-120	[-1, 512, 12, 8]	1,024
ReLU-121	[-1, 512, 12, 8]	0
BasicBlock-122	[-1, 512, 12, 8]	0
AdaptiveAvgPool2d-123	[-1, 512, 1, 1]	0
Linear-124	[-1, 8]	4,104
ResNet-125	[-1, 8]	0
Linear-126	[-1, 2]	18

=====

Total params: 21,288,794

Trainable params: 4,122

Non-trainable params: 21,284,672

Input size (MB): 1.07

Forward/backward pass size (MB): 183.55

Params size (MB): 81.21

Estimated Total Size (MB): 265.83
